

Свёрточные нейронные сети: CIFAR10

Выполнила: Трофимова Екатерина Александровна, 20223

Теория CNN

В этом ноутбке мы посмотрим, насколько хорошо **CNN** будут предсказывать классы на более сложном датасете картинок -- **CIFAR10**.

Внимание: Рассматривается *задача классификации изображений*.

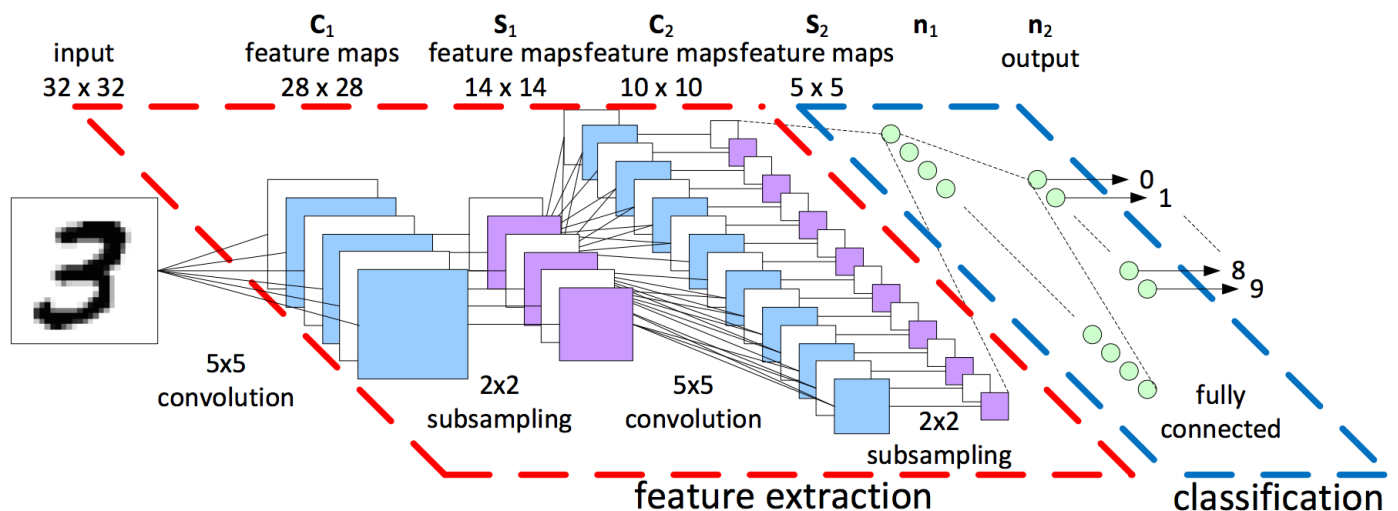
Свёрточная нейросеть (Convolutional Neural Network, CNN) - это многослойная нейросеть, имеющая в своей архитектуре помимо *полносвязных слоёв* (а иногда их может и не быть) ещё и **свёрточные слои (Conv Layers)** и **pooling-слои (Pool Layers)**.

Собственно, название такое эти сети получили потому, что в основе их работы лежит операция **свёртки**.

Сразу же стоит сказать, что свёрточные нейросети **были придуманы прежде всего для задач, связанных с изображениями**, следовательно, на вход они тоже "ожидают" изображение.

- Например, вот так выглядит неглубокая свёрточная нейросеть, имеющая такую архитектуру:

Input -> Conv 5x5 -> Pool 2x2 -> Conv 5x5 -> Pool 2x2 -> FC -> Output



Свёрточные нейросети (простые, есть и намного более продвинутые) почти всегда строятся по следующему правилу:

INPUT -> [[CONV -> RELU]*N -> POOL?]*M -> [FC -> RELU]*L -> FC

то есть:

1). Входной слой: batch картинок -- тензор размера $(batch_size, H, W, C)$ или $(batch_size, C, H, W)$

2). M блоков ($M \geq 0$) из свёрток и **pooling**-ов, причём именно в том порядке, как в формуле выше. Все эти M блоков вместе называют **feature extractor** свёрточной нейросети, потому что эта часть сети отвечает непосредственно за формирование новых, более сложных признаков поверх тех, которые подаются (то есть, по аналогии с **MLP** мы опять же переходим к новому признаковому пространству, однако здесь оно строится

аналогии с $1D$, мы опять же переходим к новому признаковому пространству, однако здесь это строится сложнее, чем в обычных многослойных сетях, поскольку используется операция свёртки)

3). L штук **FullyConnected**-слоёв (с активациями). Эту часть из L **FC**-слоёв называют **classifier**, поскольку эти слои отвечают непосредственно за предсказание нужного класса (сейчас рассматривается задача классификации изображений).

Свёрточная нейросеть на PyTorch

Ещё раз напомним про основные компоненты нейросети:

- непосредственно, сама **архитектура** нейросети (сюда входят типы функций активации у каждого нейрона);
- начальная **инициализация** весов каждого слоя;
- метод **оптимизации** нейросети (сюда ещё входит метод изменения `learning_rate`);
- размер **батчей** (`batch_size`);
- количество **эпох** обучения (`num_epochs`);
- **функция потерь** (`loss`);
- тип **регуляризации** нейросети (`weight_decay`, для каждого слоя можно свой);


То, что связано с *данными и задачей*:

- само **качество** выборки (непротиворечивость, чистота, корректность постановки задачи);
- **размер** выборки;

Так как мы сейчас рассматриваем **архитектуру CNN**, то, помимо этих компонент, в свёрточной нейросети можно настроить следующие вещи:

- (в каждом **ConvLayer**) размер фильтров (окна свёртки) (`kernel_size`)
- (в каждом **ConvLayer**) количество фильтров (`out_channels`)
- (в каждом **ConvLayer**) размер шага окна свёртки (**stride**) (`stride`)
- (в каждом **ConvLayer**) тип **padding'a** (`padding`)
- (в каждом **PoolLayer**) размер окна **pooling'a** (`kernel_size`)
- (в каждом **PoolLayer**) шаг окна **pooling'a** (`stride`)
- (в каждом **PoolLayer**) тип **pooling'a** (`pool_type`)
- (в каждом **PoolLayer**) тип **padding'a** (`padding`)

CIFAR10



CIFAR10: это набор из **60k** картинок **32x32x3**, **50k** которых составляют обучающую выборку, и оставшиеся **10k** - тестовую. Классов в этом датасете **10**: 'plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'.

In [2]:

```
# !pip install torch torchvision
```

In [3]:

```
import torch
import torchvision
from torchvision import transforms
from tqdm import tqdm_notebook

import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

In [4]:

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='../pytorch_data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='../pytorch_data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ../pytorch_data/cifar-10-python.tar.gz

Extracting ../pytorch_data/cifar-10-python.tar.gz to ../pytorch_data
Files already downloaded and verified

In []:

```
trainset.data
```

In [6]:

```
trainloader.dataset.train_list[0]
```


Out[6]:


```
['data_batch_1', 'c99cafc152244af753f735de768cd75f']
```

In [7]:

```
# случайный индекс от 0 до размера тренировочной выборки
i = np.random.randint(low=0, high=50000)

plt.imshow(trainloader.dataset.data[i]);
```





CNN для предсказания на CIFAR10.

Напишем свёрточную нейросеть для предсказания на **CIFAR10**

In [8]:

```
# Подключение зависимостей

import torch.nn as nn
import torch.nn.functional as F
from tqdm import tqdm_notebook
from torch.optim import lr_scheduler
```

Вспомогательные функции

In [9]:

```
# Попытка ускорить вычисления за счет gpu

def get_default_device():
    """Pick GPU if available, else CPU"""
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device)

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)
```

In [10]:

```
device = get_default_device()
device
```

Out[10]:

```
device(type='cpu')
```

In [11]:

```
#trainloader = DeviceDataLoader(trainloader, device)
```

```
#testloader = DeviceDataLoader(testloader, device)
```

In [12]:

```
# Функция для обучения модели

def train(net, epoch_num = 5, learning_rate = 1e-3):

    loss_fn = torch.nn.CrossEntropyLoss()

    optimizer = torch.optim.Adam(net.parameters(), lr=learning_rate)
    # динамически изменяем LR
    scheduler = lr_scheduler.CosineAnnealingLR(optimizer, T_max=epoch_num)

    # итерируемся
    for epoch in tqdm_notebook(range(epoch_num)):

        scheduler.step()

        running_loss = 0.0
        for i, batch in enumerate(tqdm_notebook(trainloader)):
            # так получаем текущий батч
            X_batch, y_batch = batch

            # обнуляем веса
            optimizer.zero_grad()

            # forward + backward + optimize
            y_pred = net(X_batch)
            loss = loss_fn(y_pred, y_batch)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            # выводим качество каждые 2000 батчей
            if i % 2000 == 1999:
                print('[%d, %5d] loss: %.3f' %
                      (epoch + 1, i + 1, running_loss / 2000))
                running_loss = 0.0

    print('Обучение закончено')
```

In [13]:

```
# Функция для проверки качества

def check_accuracy(net):
    class_correct = list(0. for i in range(10))
    class_total = list(0. for i in range(10))

    with torch.no_grad():
        for data in testloader:
            images, labels = data
            y_pred = net(images)
            _, predicted = torch.max(y_pred, 1)
            c = (predicted == labels).squeeze()
            for i in range(4):
                label = labels[i]
                class_correct[label] += c[i].item()
                class_total[label] += 1

    avg_accuracy = 0

    for i in range(10):
        print('Accuracy of %5s : %2d %%' % (classes[i], 100 * class_correct[i] / class_total[i]))
        avg_accuracy += 100 * class_correct[i] / class_total[i]

    print('Avg accuracy %2d %%' % (avg_accuracy / 10))
```

In [14]:

```
# Функция для визуальной проверки результата

def visualize_result(net, index):
    image = testloader.dataset.data[index]
    plt.imshow(image)

    image = transform(image)  # не забудем отмасштабировать!

    y_pred = net(image.view(1, 3, 32, 32))
    _, predicted = torch.max(y_pred, 1)

    plt.title(f'Predicted: {classes[predicted.numpy()[0]]}')
```

Базовая архитектура

In [58]:

```
class SimpleConvNet(torch.nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(SimpleConvNet, self).__init__()
        # feature extractor
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=6, kernel_size=5)
        self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)
        # classifier
        self.fc1 = nn.Linear(5 * 5 * 16, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 5 * 5 * 16)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

In [59]:

```
net = SimpleConvNet()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate
```

```
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.893
[1, 4000] loss: 1.628
[1, 6000] loss: 1.506
[1, 8000] loss: 1.471
[1, 10000] loss: 1.431
[1, 12000] loss: 1.359
```

```
[2, 2000] loss: 1.288
```

```
[2, 2000] loss: 1.289
[2, 4000] loss: 1.288
[2, 6000] loss: 1.268
[2, 8000] loss: 1.272
[2, 10000] loss: 1.248
[2, 12000] loss: 1.236
```

```
[3, 2000] loss: 1.152
[3, 4000] loss: 1.146
[3, 6000] loss: 1.132
[3, 8000] loss: 1.153
[3, 10000] loss: 1.139
[3, 12000] loss: 1.120
```

```
[4, 2000] loss: 1.031
[4, 4000] loss: 1.037
[4, 6000] loss: 1.026
[4, 8000] loss: 1.052
[4, 10000] loss: 1.049
[4, 12000] loss: 1.039
```

```
[5, 2000] loss: 0.947
[5, 4000] loss: 0.950
[5, 6000] loss: 0.949
[5, 8000] loss: 0.938
[5, 10000] loss: 0.947
[5, 12000] loss: 0.952
```

```
[6, 2000] loss: 0.837
[6, 4000] loss: 0.845
[6, 6000] loss: 0.862
[6, 8000] loss: 0.884
[6, 10000] loss: 0.867
[6, 12000] loss: 0.878
```

```
[7, 2000] loss: 0.798
[7, 4000] loss: 0.780
[7, 6000] loss: 0.796
[7, 8000] loss: 0.778
[7, 10000] loss: 0.776
[7, 12000] loss: 0.806
```

```
[8, 2000] loss: 0.717
[8, 4000] loss: 0.752
[8, 6000] loss: 0.720
[8, 8000] loss: 0.727
[8, 10000] loss: 0.732
[8, 12000] loss: 0.724
```

```
[9, 2000] loss: 0.671
[9, 4000] loss: 0.709
[9, 6000] loss: 0.685
[9, 8000] loss: 0.682
[9, 10000] loss: 0.712
[9, 12000] loss: 0.696
```

```
[10, 2000] loss: 0.699
[10, 4000] loss: 0.677
[10, 6000] loss: 0.690
[10, 8000] loss: 0.686
[10, 10000] loss: 0.661
[10, 12000] loss: 0.669
```

Обучение закончено

Посмотрим на **accuracy** на тестовом датасете:

In [60]:

```
check_accuracy(net)
```


```
Accuracy of plane : 70 %  
Accuracy of car : 77 %  
Accuracy of bird : 50 %  
Accuracy of cat : 44 %  
Accuracy of deer : 58 %  
Accuracy of dog : 51 %  
Accuracy of frog : 74 %  
Accuracy of horse : 71 %  
Accuracy of ship : 77 %  
Accuracy of truck : 72 %  
Avg accuracy 64 %
```

При базовой архитектуре наблюдается средняя точность в районе **64%** на **10** эпохах. Минимальная точность класса **44%**. Среднее время вычисления на эпоху = **1:10**

Проверим работу нейросети визуально (позапускайте ячейку несколько раз):

In [37]:

```
i = np.random.randint(low=0, high=10000)  
visualize_result(net, i)
```



Эксперименты с числом сверточных слоев и каналов

Попробуем просто добавить новый сверточный слой

In [68]:

```
class ConvNet_3CL(nn.Module):  
    def __init__(self):  
        # вызов конструктора класса nn.Module()  
        super(ConvNet_3CL, self).__init__()  
  
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)  
  
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=6, kernel_size=5)  
        self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)  
        self.conv3 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=5)  
  
        self.fc1 = nn.Linear(3 * 3 * 32, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.relu(self.conv1(x))  
        x = self.pool(x)  
        x = F.relu(self.conv2(x))  
        x = F.relu(self.conv3(x))  
        x = self.pool(x)
```



```

x = x.view(-1, 3 * 3 * 32)
x = F.relu(self.fc1(x))
x = F.relu(self.fc2(x))
x = self.fc3(x)
return x

```

In [69]:

```

net = ConvNet_3CL()
train(net, learning_rate=0.001, epoch_num=10)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
  if sys.path[0] == '':

```

```

/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate

```

```

  "https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`

```

```

[1, 2000] loss: 2.017
[1, 4000] loss: 1.828
[1, 6000] loss: 1.682
[1, 8000] loss: 1.572
[1, 10000] loss: 1.531
[1, 12000] loss: 1.507

```

```

[2, 2000] loss: 1.458
[2, 4000] loss: 1.427
[2, 6000] loss: 1.412
[2, 8000] loss: 1.393
[2, 10000] loss: 1.382
[2, 12000] loss: 1.362

```

```

[3, 2000] loss: 1.302
[3, 4000] loss: 1.314
[3, 6000] loss: 1.281
[3, 8000] loss: 1.275
[3, 10000] loss: 1.275
[3, 12000] loss: 1.268

```

```

[4, 2000] loss: 1.191
[4, 4000] loss: 1.183
[4, 6000] loss: 1.194
[4, 8000] loss: 1.202
[4, 10000] loss: 1.178
[4, 12000] loss: 1.187

```

```

[5, 2000] loss: 1.083
[5, 4000] loss: 1.109
[5, 6000] loss: 1.101
[5, 8000] loss: 1.110
[5, 10000] loss: 1.115
[5, 12000] loss: 1.095

```

```

[6, 2000] loss: 1.014
[6, 4000] loss: 1.024
[6, 6000] loss: 1.003
[6, 8000] loss: 1.015
[6, 10000] loss: 1.026
[6, 12000] loss: 1.029

```

```
[7, 2000] loss: 0.938
[7, 4000] loss: 0.932
[7, 6000] loss: 0.966
[7, 8000] loss: 0.935
[7, 10000] loss: 0.929
[7, 12000] loss: 0.946
```

```
[8, 2000] loss: 0.861
[8, 4000] loss: 0.875
[8, 6000] loss: 0.898
[8, 8000] loss: 0.867
[8, 10000] loss: 0.883
[8, 12000] loss: 0.875
```

```
[9, 2000] loss: 0.829
[9, 4000] loss: 0.839
[9, 6000] loss: 0.834
[9, 8000] loss: 0.828
[9, 10000] loss: 0.844
[9, 12000] loss: 0.833
```

```
[10, 2000] loss: 0.800
[10, 4000] loss: 0.829
[10, 6000] loss: 0.820
[10, 8000] loss: 0.806
[10, 10000] loss: 0.832
[10, 12000] loss: 0.839
```

Обучение закончено

In [71]:

```
check_accuracy(net)
```

```
Accuracy of plane : 64 %
Accuracy of car : 75 %
Accuracy of bird : 44 %
Accuracy of cat : 44 %
Accuracy of deer : 54 %
Accuracy of dog : 46 %
Accuracy of frog : 72 %
Accuracy of horse : 66 %
Accuracy of ship : 75 %
Accuracy of truck : 70 %
Avg accuracy 61 %
```

Добавление еще одного сверточного слоя с малым количеством каналов отрицательно сказалось на качестве обучения (средний результат ухудшился до **61%**) и времени обучения. Примерно на **9** эпохе процесс обучения застыл. Попробуем теперь вернуться к **2м** сверточным слоям, но увеличим число каналов

In [80]:

```
class ConvNet_2Cl(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_2Cl, self).__init__()

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=5)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=5)

        self.fc1 = nn.Linear(5 * 5 * 128, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
```

```

x = F.relu(self.conv2(x))
x = self.pool(x)
x = x.view(-1, 5 * 5 * 128)
x = F.relu(self.fc1(x))
x = F.relu(self.fc2(x))
x = self.fc3(x)
return x

```

In [81]:

```

net = ConvNet_2Cl()
train(net, learning_rate=0.001, epoch_num=10)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
  if sys.path[0] == '':

```

```

/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate
  "https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`

```

```

[1, 2000] loss: 1.861
[1, 4000] loss: 1.567
[1, 6000] loss: 1.458
[1, 8000] loss: 1.388
[1, 10000] loss: 1.348
[1, 12000] loss: 1.317

```

```

[2, 2000] loss: 1.222
[2, 4000] loss: 1.187
[2, 6000] loss: 1.164
[2, 8000] loss: 1.159
[2, 10000] loss: 1.151
[2, 12000] loss: 1.104

```

```

[3, 2000] loss: 1.024
[3, 4000] loss: 1.002
[3, 6000] loss: 1.007
[3, 8000] loss: 0.995
[3, 10000] loss: 1.008
[3, 12000] loss: 1.017

```

```

[4, 2000] loss: 0.864
[4, 4000] loss: 0.889
[4, 6000] loss: 0.897
[4, 8000] loss: 0.899
[4, 10000] loss: 0.908
[4, 12000] loss: 0.894

```

```

[5, 2000] loss: 0.777
[5, 4000] loss: 0.773
[5, 6000] loss: 0.783
[5, 8000] loss: 0.784
[5, 10000] loss: 0.789
[5, 12000] loss: 0.785

```

```

[6, 2000] loss: 0.658
[6, 4000] loss: 0.674
[6, 6000] loss: 0.672
[6, 8000] loss: 0.669
[6, 10000] loss: 0.690

```

```
[6, 12000] loss: 0.668
```

```
[7, 2000] loss: 0.566
[7, 4000] loss: 0.575
[7, 6000] loss: 0.579
[7, 8000] loss: 0.581
[7, 10000] loss: 0.569
[7, 12000] loss: 0.551
```

```
[8, 2000] loss: 0.475
[8, 4000] loss: 0.489
[8, 6000] loss: 0.466
[8, 8000] loss: 0.501
[8, 10000] loss: 0.490
[8, 12000] loss: 0.495
```

```
[9, 2000] loss: 0.431
[9, 4000] loss: 0.418
[9, 6000] loss: 0.420
[9, 8000] loss: 0.431
[9, 10000] loss: 0.427
[9, 12000] loss: 0.436
```

```
[10, 2000] loss: 0.400
[10, 4000] loss: 0.420
[10, 6000] loss: 0.408
[10, 8000] loss: 0.403
[10, 10000] loss: 0.405
[10, 12000] loss: 0.416
```

Обучение закончено

In [82]:

```
check_accuracy(net)
```

```
Accuracy of plane : 76 %
Accuracy of car : 83 %
Accuracy of bird : 59 %
Accuracy of cat : 51 %
Accuracy of deer : 64 %
Accuracy of dog : 61 %
Accuracy of frog : 81 %
Accuracy of horse : 75 %
Accuracy of ship : 82 %
Accuracy of truck : 82 %
Avg accuracy 71 %
```

Увеличение числа каналов положительно сказалось на средней точности - **71%** против базовых **64%**. Но обучение в рамках эпохи теперь идет гораздо дольше. Попробуем одновременно увеличить число сверточных слоев и число каналов

In [24]:

```
class ConvNet_3CL_CH(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3CL_CH, self).__init__()

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=5)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=5)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=5)

        self.fc1 = nn.Linear(256, 120) # 1 x 1 x 256
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

```
def forward(self, x):
    x = F.relu(self.conv1(x))
    x = self.pool(x)
    x = F.relu(self.conv2(x))
    x = self.pool(x)
    x = F.relu(self.conv3(x))
    x = x.view(-1, 256)
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x
```

In [25]:

```
net = ConvNet_3CL_CH()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate
```

```
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.913
[1, 4000] loss: 1.647
[1, 6000] loss: 1.544
[1, 8000] loss: 1.488
[1, 10000] loss: 1.459
[1, 12000] loss: 1.390
```

```
[2, 2000] loss: 1.330
[2, 4000] loss: 1.299
[2, 6000] loss: 1.279
[2, 8000] loss: 1.257
[2, 10000] loss: 1.230
[2, 12000] loss: 1.227
```

```
[3, 2000] loss: 1.139
[3, 4000] loss: 1.139
[3, 6000] loss: 1.127
[3, 8000] loss: 1.101
[3, 10000] loss: 1.111
[3, 12000] loss: 1.075
```

```
[4, 2000] loss: 0.996
[4, 4000] loss: 1.012
[4, 6000] loss: 1.014
[4, 8000] loss: 0.997
[4, 10000] loss: 0.991
[4, 12000] loss: 0.973
```

```
[5, 2000] loss: 0.875
[5, 4000] loss: 0.875
[5, 6000] loss: 0.890
[5, 8000] loss: 0.896
[5, 10000] loss: 0.892
[5, 12000] loss: 0.902
```

```
[6, 2000] loss: 0.777
```

```
[6, 4000] loss: 0.783
[6, 6000] loss: 0.794
[6, 8000] loss: 0.776
[6, 10000] loss: 0.798
[6, 12000] loss: 0.769
```

```
[7, 2000] loss: 0.678
[7, 4000] loss: 0.694
[7, 6000] loss: 0.672
[7, 8000] loss: 0.679
[7, 10000] loss: 0.680
[7, 12000] loss: 0.703
```

```
[8, 2000] loss: 0.607
[8, 4000] loss: 0.589
[8, 6000] loss: 0.600
[8, 8000] loss: 0.596
[8, 10000] loss: 0.588
[8, 12000] loss: 0.591
```

```
[9, 2000] loss: 0.529
[9, 4000] loss: 0.520
[9, 6000] loss: 0.524
[9, 8000] loss: 0.544
[9, 10000] loss: 0.561
[9, 12000] loss: 0.528
```

```
[10, 2000] loss: 0.506
[10, 4000] loss: 0.518
[10, 6000] loss: 0.514
[10, 8000] loss: 0.518
[10, 10000] loss: 0.534
[10, 12000] loss: 0.513
```

Обучение закончено

In [26]:

```
check_accuracy(net)
```

```
Accuracy of plane : 75 %
Accuracy of   car : 81 %
Accuracy of  bird : 54 %
Accuracy of   cat : 49 %
Accuracy of  deer : 67 %
Accuracy of   dog : 59 %
Accuracy of  frog : 79 %
Accuracy of horse : 73 %
Accuracy of  ship : 81 %
Accuracy of truck : 78 %
Avg accuracy 70 %
```

Средняя точность стала чуть ниже - **70%** против **71%** на двух слоях. При этом сильно возросло время обучения. Теперь попробуем изменить размер ядра свертки для случая с двумя слоями

In [17]:

```
class ConvNet_2Cl_3KS(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_2Cl_3KS, self).__init__()

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)

        self.fc1 = nn.Linear(6 * 6 * 128, 120)
        self.fc2 = nn.Linear(120, 84)
```

```

        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        x = x.view(-1, 6 * 6 * 128)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

```

In [18]:

```

net = ConvNet_2Cl_3KS()
train(net, learning_rate=0.001, epoch_num=10)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':

```

```

/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate

```

```

"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`

```

```

[1, 2000] loss: 1.801
[1, 4000] loss: 1.463
[1, 6000] loss: 1.395
[1, 8000] loss: 1.290
[1, 10000] loss: 1.256
[1, 12000] loss: 1.211

```

```

[2, 2000] loss: 1.151
[2, 4000] loss: 1.105
[2, 6000] loss: 1.077
[2, 8000] loss: 1.052
[2, 10000] loss: 1.047
[2, 12000] loss: 1.039

```

```

[3, 2000] loss: 0.930
[3, 4000] loss: 0.924
[3, 6000] loss: 0.934
[3, 8000] loss: 0.945
[3, 10000] loss: 0.940
[3, 12000] loss: 0.934

```

```

[4, 2000] loss: 0.805
[4, 4000] loss: 0.819
[4, 6000] loss: 0.839
[4, 8000] loss: 0.839
[4, 10000] loss: 0.847
[4, 12000] loss: 0.812

```

```

[5, 2000] loss: 0.714
[5, 4000] loss: 0.723
[5, 6000] loss: 0.730
[5, 8000] loss: 0.736
[5, 10000] loss: 0.749
[5, 12000] loss: 0.738

```

```
[6, 2000] loss: 0.632
[6, 4000] loss: 0.627
[6, 6000] loss: 0.637
[6, 8000] loss: 0.647
[6, 10000] loss: 0.659
[6, 12000] loss: 0.649
```

```
[7, 2000] loss: 0.557
[7, 4000] loss: 0.558
[7, 6000] loss: 0.559
[7, 8000] loss: 0.555
[7, 10000] loss: 0.575
[7, 12000] loss: 0.564
```

```
[8, 2000] loss: 0.511
[8, 4000] loss: 0.479
[8, 6000] loss: 0.505
[8, 8000] loss: 0.505
[8, 10000] loss: 0.491
[8, 12000] loss: 0.496
```

```
[9, 2000] loss: 0.465
[9, 4000] loss: 0.453
[9, 6000] loss: 0.462
[9, 8000] loss: 0.467
[9, 10000] loss: 0.473
[9, 12000] loss: 0.448
```

```
[10, 2000] loss: 0.439
[10, 4000] loss: 0.446
[10, 6000] loss: 0.455
[10, 8000] loss: 0.442
[10, 10000] loss: 0.450
[10, 12000] loss: 0.457
```

Обучение закончено

In [19]:

```
check_accuracy(net)
```

```
Accuracy of plane : 74 %
Accuracy of   car : 83 %
Accuracy of  bird : 57 %
Accuracy of   cat : 52 %
Accuracy of  deer : 65 %
Accuracy of   dog : 63 %
Accuracy of  frog : 77 %
Accuracy of horse : 77 %
Accuracy of  ship : 81 %
Accuracy of truck : 81 %
Avg accuracy 71 %
```

Для двух слоев изменение размера ядра свертки не дало существенных изменений. Теперь посмотрим на **3х** слоях

In [27]:

```
class ConvNet_3CL_CH_3KS(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3CL_CH_3KS, self).__init__()

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)
```



```

self.fc1 = nn.Linear(2 * 2 * 256, 120)
self.fc2 = nn.Linear(120, 84)
self.fc3 = nn.Linear(84, 10)

```

```

def forward(self, x):
    x = F.relu(self.conv1(x))
    x = self.pool(x)
    x = F.relu(self.conv2(x))
    x = self.pool(x)
    x = F.relu(self.conv3(x))
    x = self.pool(x)
    x = x.view(-1, 2 * 2 * 256)
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x

```

In [28]:

```

net = ConvNet_3CL_CH_3KS()
train(net, learning_rate=0.001, epoch_num=10)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
  if sys.path[0] == '':

```

```

/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate

```

```

  "https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`

```

```

[1, 2000] loss: 1.869
[1, 4000] loss: 1.537
[1, 6000] loss: 1.431
[1, 8000] loss: 1.371
[1, 10000] loss: 1.279
[1, 12000] loss: 1.251

```

```

[2, 2000] loss: 1.146
[2, 4000] loss: 1.095
[2, 6000] loss: 1.078
[2, 8000] loss: 1.065
[2, 10000] loss: 1.071
[2, 12000] loss: 1.057

```

```

[3, 2000] loss: 0.934
[3, 4000] loss: 0.945
[3, 6000] loss: 0.943
[3, 8000] loss: 0.919
[3, 10000] loss: 0.924
[3, 12000] loss: 0.914

```

```

[4, 2000] loss: 0.791
[4, 4000] loss: 0.812
[4, 6000] loss: 0.798
[4, 8000] loss: 0.824
[4, 10000] loss: 0.808
[4, 12000] loss: 0.813

```

```

[5, 2000] loss: 0.692
[5, 4000] loss: 0.697
[5, 6000] loss: 0.693

```

```
[5, 8000] loss: 0.722
[5, 10000] loss: 0.702
[5, 12000] loss: 0.689
```

```
[6, 2000] loss: 0.589
[6, 4000] loss: 0.567
[6, 6000] loss: 0.565
[6, 8000] loss: 0.579
[6, 10000] loss: 0.593
[6, 12000] loss: 0.611
```

```
[7, 2000] loss: 0.471
[7, 4000] loss: 0.479
[7, 6000] loss: 0.489
[7, 8000] loss: 0.469
[7, 10000] loss: 0.478
[7, 12000] loss: 0.490
```

```
[8, 2000] loss: 0.375
[8, 4000] loss: 0.401
[8, 6000] loss: 0.403
[8, 8000] loss: 0.379
[8, 10000] loss: 0.406
[8, 12000] loss: 0.399
```

```
[9, 2000] loss: 0.351
[9, 4000] loss: 0.339
[9, 6000] loss: 0.323
[9, 8000] loss: 0.336
[9, 10000] loss: 0.339
[9, 12000] loss: 0.343
```

```
[10, 2000] loss: 0.318
[10, 4000] loss: 0.321
[10, 6000] loss: 0.315
[10, 8000] loss: 0.324
[10, 10000] loss: 0.319
[10, 12000] loss: 0.321
```

Обучение закончено

In [29]:

```
check_accuracy(net)
```

```
Accuracy of plane : 79 %
Accuracy of car : 85 %
Accuracy of bird : 60 %
Accuracy of cat : 54 %
Accuracy of deer : 70 %
Accuracy of dog : 61 %
Accuracy of frog : 79 %
Accuracy of horse : 75 %
Accuracy of ship : 82 %
Accuracy of truck : 84 %
Avg accuracy 73 %
```

А вот на **3х** слоях уже наблюдается небольшой прирост: **73%** против **70%**. Минимальная точность возросла до **54%**.

Добавление слоев и каналов позволило обогатить пространство признаков и улучшить тем самым результат классификации

Эксперименты с пулингом и нормализацией

Теперь попробуем поменять тип пулинга с **max** на **avg**

In [20]:

```
class ConvNet_3Cl_3KS_AvgPool(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3Cl_3KS_AvgPool, self).__init__()

        self.pool = nn.AvgPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)

        self.fc1 = nn.Linear(2 * 2 * 256, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        x = F.relu(self.conv3(x))
        x = self.pool(x)
        x = x.view(-1, 2 * 2 * 256)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

In [21]:

```
net = ConvNet_3Cl_3KS_AvgPool()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.924
[1, 4000] loss: 1.625
[1, 6000] loss: 1.478
[1, 8000] loss: 1.390
[1, 10000] loss: 1.328
[1, 12000] loss: 1.277
```

```
[2, 2000] loss: 1.203
[2, 4000] loss: 1.143
[2, 6000] loss: 1.126
[2, 8000] loss: 1.073
[2, 10000] loss: 1.062
[2, 12000] loss: 1.025
```

```
[3, 2000] loss: 0.942
[3, 4000] loss: 0.919
[3, 6000] loss: 0.919
[3, 8000] loss: 0.892
[3, 10000] loss: 0.890
[3, 12000] loss: 0.898
```

```
[4, 2000] loss: 0.772
[4, 4000] loss: 0.784
[4, 6000] loss: 0.778
[4, 8000] loss: 0.781
[4, 10000] loss: 0.763
[4, 12000] loss: 0.778
```

```
[5, 2000] loss: 0.659
[5, 4000] loss: 0.666
[5, 6000] loss: 0.687
[5, 8000] loss: 0.645
[5, 10000] loss: 0.661
[5, 12000] loss: 0.657
```

```
[6, 2000] loss: 0.552
[6, 4000] loss: 0.547
[6, 6000] loss: 0.556
[6, 8000] loss: 0.566
[6, 10000] loss: 0.556
[6, 12000] loss: 0.574
```

```
[7, 2000] loss: 0.466
[7, 4000] loss: 0.458
[7, 6000] loss: 0.473
[7, 8000] loss: 0.475
[7, 10000] loss: 0.464
[7, 12000] loss: 0.472
```

```
[8, 2000] loss: 0.398
[8, 4000] loss: 0.388
[8, 6000] loss: 0.401
[8, 8000] loss: 0.375
[8, 10000] loss: 0.390
[8, 12000] loss: 0.400
```

```
[9, 2000] loss: 0.347
[9, 4000] loss: 0.355
[9, 6000] loss: 0.355
[9, 8000] loss: 0.349
[9, 10000] loss: 0.340
[9, 12000] loss: 0.338
```

```
[10, 2000] loss: 0.332
[10, 4000] loss: 0.332
[10, 6000] loss: 0.332
[10, 8000] loss: 0.330
[10, 10000] loss: 0.340
[10, 12000] loss: 0.334
```

Обучение закончено

In [22]:

```
check_accuracy(net)
```

```
Accuracy of plane : 79 %
Accuracy of   car : 84 %
Accuracy of  bird : 64 %
Accuracy of   cat : 57 %
Accuracy of  deer : 74 %
Accuracy of   dog : 64 %
Accuracy of  frog : 82 %
Accuracy of horse : 78 %
Accuracy of  ship : 85 %
Accuracy of truck : 83 %
Avg accuracy 75 %
```

Смена типа пулинга с **max** на **avg** еще немного улучшила результат: средняя точность **75%** против **73%**.

минимальная - **57%** против **54%**. Т.е. положение признака оказалось немного важнее его наличия. Попробуем добавить нормализацию

In [23]:

```
class ConvNet_3Cl_3KS_AvgPool_BN(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3Cl_3KS_AvgPool_BN, self).__init__()

        self.pool = nn.AvgPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.bn1 = nn.BatchNorm2d(64)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.bn2 = nn.BatchNorm2d(128)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)
        self.bn3 = nn.BatchNorm2d(256)

        self.fc1 = nn.Linear(2 * 2 * 256, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.bn1(F.relu(self.conv1(x)))
        x = self.pool(x)
        x = self.bn2(F.relu(self.conv2(x)))
        x = self.pool(x)
        x = self.bn3(F.relu(self.conv3(x)))
        x = self.pool(x)
        x = x.view(-1, 2 * 2 * 256)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

In [24]:

```
net = ConvNet_3Cl_3KS_AvgPool_BN()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
  if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate
```

```
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.862
[1, 4000] loss: 1.600
[1, 6000] loss: 1.479
[1, 8000] loss: 1.357
[1, 10000] loss: 1.289
[1, 12000] loss: 1.241
```

```
[2, 2000] loss: 1.110
[2, 4000] loss: 1.081
[2, 6000] loss: 1.094
[2, 8000] loss: 1.068
[2, 10000] loss: 1.066
[2, 12000] loss: 1.029
```

```
[3, 2000] loss: 0.943
[3, 4000] loss: 0.903
[3, 6000] loss: 0.941
[3, 8000] loss: 0.920
[3, 10000] loss: 0.896
[3, 12000] loss: 0.898
```

```
[4, 2000] loss: 0.770
[4, 4000] loss: 0.777
[4, 6000] loss: 0.787
[4, 8000] loss: 0.769
[4, 10000] loss: 0.794
[4, 12000] loss: 0.782
```

```
[5, 2000] loss: 0.660
[5, 4000] loss: 0.668
[5, 6000] loss: 0.651
[5, 8000] loss: 0.677
[5, 10000] loss: 0.651
[5, 12000] loss: 0.649
```

```
[6, 2000] loss: 0.547
[6, 4000] loss: 0.538
[6, 6000] loss: 0.551
[6, 8000] loss: 0.546
[6, 10000] loss: 0.535
[6, 12000] loss: 0.542
```

```
[7, 2000] loss: 0.451
[7, 4000] loss: 0.444
[7, 6000] loss: 0.451
[7, 8000] loss: 0.452
[7, 10000] loss: 0.446
[7, 12000] loss: 0.438
```

```
[8, 2000] loss: 0.377
[8, 4000] loss: 0.371
[8, 6000] loss: 0.381
[8, 8000] loss: 0.376
[8, 10000] loss: 0.374
[8, 12000] loss: 0.366
```

```
[9, 2000] loss: 0.321
[9, 4000] loss: 0.342
[9, 6000] loss: 0.330
[9, 8000] loss: 0.343
[9, 10000] loss: 0.323
[9, 12000] loss: 0.337
```

```
[10, 2000] loss: 0.322
[10, 4000] loss: 0.319
[10, 6000] loss: 0.319
[10, 8000] loss: 0.331
[10, 10000] loss: 0.325
[10, 12000] loss: 0.329
```

Обучение закончено

In [25]:

```
check_accuracy(net)
```

```
Accuracy of plane : 81 %
Accuracy of car : 85 %
Accuracy of bird : 65 %
Accuracy of cat : 57 %
Accuracy of deer : 73 %
Accuracy of dog : 64 %
```

```
Accuracy of frog : 81 %
Accuracy of horse : 81 %
Accuracy of ship : 83 %
Accuracy of truck : 82 %
Avg accuracy 75 %
```

Нормализация не повлияла на результат

Эксперимент с функцией активации

Обычно в качестве функции активации сверточных слоев используют функцию **ReLU**. Рассматривать функции активации вроде сигмоидной или тангенциальной мы не будем, т.к. они приводят к проблемам с затуханием или увеличением градиентов. Вместо этого попробуем использовать **ELU**, которая сохраняет преимущества **ReLU** и помогает избежать проблемы умирающего **ReLU**

In [27]:

```
class ConvNet_3Cl_3KS_AvgPool_ELU(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3Cl_3KS_AvgPool_ELU, self).__init__()

        self.pool = nn.AvgPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)

        self.fc1 = nn.Linear(2 * 2 * 256, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.elu(self.conv1(x))
        x = self.pool(x)
        x = F.elu(self.conv2(x))
        x = self.pool(x)
        x = F.elu(self.conv3(x))
        x = self.pool(x)
        x = x.view(-1, 2 * 2 * 256)
        x = F.elu(self.fc1(x))
        x = F.elu(self.fc2(x))
        x = self.fc3(x)
        return x
```

In [30]:

```
net = ConvNet_3Cl_3KS_AvgPool_ELU()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate
```

```
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.932
[1. 4000] loss: 1.686
```

[1, 2000] loss: 1.577
[1, 4000] loss: 1.484
[1, 6000] loss: 1.433
[1, 8000] loss: 1.373

[2, 2000] loss: 1.255
[2, 4000] loss: 1.227
[2, 6000] loss: 1.192
[2, 8000] loss: 1.178
[2, 10000] loss: 1.161
[2, 12000] loss: 1.134

[3, 2000] loss: 0.965
[3, 4000] loss: 0.994
[3, 6000] loss: 0.978
[3, 8000] loss: 0.978
[3, 10000] loss: 0.967
[3, 12000] loss: 0.973

[4, 2000] loss: 0.803
[4, 4000] loss: 0.784
[4, 6000] loss: 0.800
[4, 8000] loss: 0.815
[4, 10000] loss: 0.808
[4, 12000] loss: 0.828

[5, 2000] loss: 0.639
[5, 4000] loss: 0.634
[5, 6000] loss: 0.651
[5, 8000] loss: 0.655
[5, 10000] loss: 0.642
[5, 12000] loss: 0.645

[6, 2000] loss: 0.480
[6, 4000] loss: 0.465
[6, 6000] loss: 0.501
[6, 8000] loss: 0.501
[6, 10000] loss: 0.508
[6, 12000] loss: 0.495

[7, 2000] loss: 0.340
[7, 4000] loss: 0.340
[7, 6000] loss: 0.340
[7, 8000] loss: 0.358
[7, 10000] loss: 0.359
[7, 12000] loss: 0.369

[8, 2000] loss: 0.235
[8, 4000] loss: 0.242
[8, 6000] loss: 0.251
[8, 8000] loss: 0.240
[8, 10000] loss: 0.253
[8, 12000] loss: 0.231

[9, 2000] loss: 0.190
[9, 4000] loss: 0.170
[9, 6000] loss: 0.178
[9, 8000] loss: 0.174
[9, 10000] loss: 0.180
[9, 12000] loss: 0.181

[10, 2000] loss: 0.163
[10, 4000] loss: 0.160
[10, 6000] loss: 0.145
[10, 8000] loss: 0.164
[10, 10000] loss: 0.156
[10, 12000] loss: 0.162

обучение завершено

Обучение закончено

In [46]:

```
check_accuracy(net)
```

```
Accuracy of plane : 78 %
Accuracy of   car : 83 %
Accuracy of  bird : 63 %
Accuracy of   cat : 55 %
Accuracy of  deer : 66 %
Accuracy of   dog : 64 %
Accuracy of  frog : 78 %
Accuracy of horse : 74 %
Accuracy of  ship : 81 %
Accuracy of truck : 80 %
Avg accuracy 72 %
```

Несмотря на то, что значение функции потерь теперь меньше, точность тоже упала. **ELU** не дает нам выигрыша на текущей архитектуре

Эксперимент с числом полносвязных слоев

Начнем с простого - уберем 1 слой

In [14]:

```
class ConvNet_3C1_3KS_AvgPool_2F1(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(ConvNet_3C1_3KS_AvgPool_2F1, self).__init__()

        self.pool = nn.AvgPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)

        self.fc1 = nn.Linear(2 * 2 * 256, 512)
        self.fc3 = nn.Linear(512, 10)

    def forward(self, x):
        x = F.elu(self.conv1(x))
        x = self.pool(x)
        x = F.elu(self.conv2(x))
        x = self.pool(x)
        x = F.elu(self.conv3(x))
        x = self.pool(x)
        x = x.view(-1, 2 * 2 * 256)
        x = F.elu(self.fc1(x))
        x = self.fc3(x)
        return x
```

In [15]:

```
net = ConvNet_3C1_3KS_AvgPool_2F1()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
  if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjus
t-learning-rate
```

```
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 2.004
[1, 4000] loss: 1.805
[1, 6000] loss: 1.700
[1, 8000] loss: 1.636
[1, 10000] loss: 1.590
[1, 12000] loss: 1.529
```

```
[2, 2000] loss: 1.410
[2, 4000] loss: 1.411
[2, 6000] loss: 1.376
[2, 8000] loss: 1.355
[2, 10000] loss: 1.306
[2, 12000] loss: 1.327
```

```
[3, 2000] loss: 1.140
[3, 4000] loss: 1.174
[3, 6000] loss: 1.159
[3, 8000] loss: 1.146
[3, 10000] loss: 1.167
[3, 12000] loss: 1.126
```

```
[4, 2000] loss: 0.886
[4, 4000] loss: 0.953
[4, 6000] loss: 0.900
[4, 8000] loss: 0.969
[4, 10000] loss: 0.936
[4, 12000] loss: 0.958
```

```
[5, 2000] loss: 0.654
[5, 4000] loss: 0.692
[5, 6000] loss: 0.686
[5, 8000] loss: 0.695
[5, 10000] loss: 0.737
[5, 12000] loss: 0.715
```

```
[6, 2000] loss: 0.439
[6, 4000] loss: 0.430
[6, 6000] loss: 0.449
[6, 8000] loss: 0.449
[6, 10000] loss: 0.460
[6, 12000] loss: 0.463
```

```
[7, 2000] loss: 0.232
[7, 4000] loss: 0.229
[7, 6000] loss: 0.235
[7, 8000] loss: 0.247
[7, 10000] loss: 0.247
[7, 12000] loss: 0.247
```

```
[8, 2000] loss: 0.109
[8, 4000] loss: 0.108
[8, 6000] loss: 0.104
[8, 8000] loss: 0.095
[8, 10000] loss: 0.104
[8, 12000] loss: 0.110
```

```
[9, 2000] loss: 0.050
[9, 4000] loss: 0.048
[9, 6000] loss: 0.049
[9, 8000] loss: 0.049
[9, 10000] loss: 0.050
[9, 12000] loss: 0.044
```

```
[10, 2000] loss: 0.036
[10, 4000] loss: 0.035
[10, 6000] loss: 0.035
[10, 8000] loss: 0.034
[10, 10000] loss: 0.029
[10, 12000] loss: 0.035
Обучение закончено
```

In [16]:

```
check_accuracy(net)
```

```
Accuracy of plane : 74 %
Accuracy of car : 80 %
Accuracy of bird : 62 %
Accuracy of cat : 53 %
Accuracy of deer : 66 %
Accuracy of dog : 59 %
Accuracy of frog : 76 %
Accuracy of horse : 76 %
Accuracy of ship : 82 %
Accuracy of truck : 79 %
Avg accuracy 70 %
```

И наоборот - добавим 1 слой

In [17]:

```
class ConvNet_3C1_3KS_AvgPool_4F1(nn.Module):
    def __init__(self):
        # ВЫЗОВ КОНСТРУКТОРА КЛАССА nn.Module()
        super(ConvNet_3C1_3KS_AvgPool_4F1, self).__init__()

        self.pool = nn.AvgPool2d(kernel_size=2, stride=2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3)
        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3)

        self.fc1 = nn.Linear(2 * 2 * 256, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 128)
        self.fc4 = nn.Linear(128, 10)

    def forward(self, x):
        x = F.elu(self.conv1(x))
        x = self.pool(x)
        x = F.elu(self.conv2(x))
        x = self.pool(x)
        x = F.elu(self.conv3(x))
        x = self.pool(x)
        x = x.view(-1, 2 * 2 * 256)
        x = F.elu(self.fc1(x))
        x = F.elu(self.fc2(x))
        x = F.elu(self.fc3(x))
        x = self.fc4(x)
        return x
```

In [18]:

```
net = ConvNet_3C1_3KS_AvgPool_4F1()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
```

you should call them in the opposite order: ``optimizer.step()`` before ``lr_scheduler.step()``. Failure to do this will result in PyTorch skipping the first value of the learning rate schedule. See more details at <https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>

"<https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>", UserWarning)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning: This function will be removed in tqdm==5.0.0
Please use ``tqdm.notebook.tqdm`` instead of ``tqdm.tqdm_notebook``

```
[1, 2000] loss: 2.077  
[1, 4000] loss: 1.844  
[1, 6000] loss: 1.739  
[1, 8000] loss: 1.624  
[1, 10000] loss: 1.540  
[1, 12000] loss: 1.471
```

```
[2, 2000] loss: 1.806  
[2, 4000] loss: 1.525  
[2, 6000] loss: 1.382  
[2, 8000] loss: 1.323  
[2, 10000] loss: 1.313  
[2, 12000] loss: 1.270
```

```
[3, 2000] loss: 1.257  
[3, 4000] loss: 1.153  
[3, 6000] loss: 1.158  
[3, 8000] loss: 1.152  
[3, 10000] loss: 1.090  
[3, 12000] loss: 1.101
```

```
[4, 2000] loss: 1.003  
[4, 4000] loss: 0.963  
[4, 6000] loss: 0.961  
[4, 8000] loss: 0.964  
[4, 10000] loss: 1.006  
[4, 12000] loss: 0.934
```

```
[5, 2000] loss: 0.785  
[5, 4000] loss: 0.841  
[5, 6000] loss: 0.800  
[5, 8000] loss: 0.804  
[5, 10000] loss: 0.811  
[5, 12000] loss: 0.808
```

```
[6, 2000] loss: 0.641  
[6, 4000] loss: 0.629  
[6, 6000] loss: 0.633  
[6, 8000] loss: 0.634  
[6, 10000] loss: 0.647  
[6, 12000] loss: 0.668
```

```
[7, 2000] loss: 0.515  
[7, 4000] loss: 0.485  
[7, 6000] loss: 0.472  
[7, 8000] loss: 0.474  
[7, 10000] loss: 0.471  
[7, 12000] loss: 0.474
```

```
[8, 2000] loss: 0.355  
[8, 4000] loss: 0.351  
[8, 6000] loss: 0.352  
[8, 8000] loss: 0.343  
[8, 10000] loss: 0.360  
[8, 12000] loss: 0.341
```

```
[9, 2000] loss: 0.265  
[9, 4000] loss: 0.266  
[9, 6000] loss: 0.265
```

```
[9, 8000] loss: 0.260
[9, 10000] loss: 0.276
[9, 12000] loss: 0.268
```

```
[10, 2000] loss: 0.242
[10, 4000] loss: 0.244
[10, 6000] loss: 0.236
[10, 8000] loss: 0.237
[10, 10000] loss: 0.239
[10, 12000] loss: 0.244
```

Обучение закончено

In [19]:

```
check_accuracy(net)
```

```
Accuracy of plane : 78 %
Accuracy of car : 82 %
Accuracy of bird : 60 %
Accuracy of cat : 56 %
Accuracy of deer : 68 %
Accuracy of dog : 55 %
Accuracy of frog : 79 %
Accuracy of horse : 74 %
Accuracy of ship : 80 %
Accuracy of truck : 79 %
Avg accuracy 71 %
```

В обоих случаях точность классификации снизилась. Для текущей архитектуры оптимальным является наличие **3** линейных слоев

Сильная архитектура, которая уже была (!) в ноутбуке

Попробуем обучить ещё более сильную нейросеть:

In [15]:

```
class StrongConvNet(nn.Module):
    def __init__(self):
        # вызов конструктора класса nn.Module()
        super(StrongConvNet, self).__init__()

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.dropout = nn.Dropout(p=0.2)

        self.conv1 = nn.Conv2d(in_channels=3, out_channels=8, kernel_size=5)
        self.bn1 = nn.BatchNorm2d(8)
        self.conv2 = nn.Conv2d(in_channels=8, out_channels=16, kernel_size=1)
        self.bn2 = nn.BatchNorm2d(16)
        self.conv3 = nn.Conv2d(in_channels=16, out_channels=16, kernel_size=3)
        self.bn3 = nn.BatchNorm2d(16)
        self.conv4 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=1)
        self.bn4 = nn.BatchNorm2d(32)
        self.conv5 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3)
        self.bn5 = nn.BatchNorm2d(32)

        self.fc1 = nn.Linear(4 * 4 * 32, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.bn1(F.relu(self.conv1(x)))
        x = self.pool(x)
        x = self.bn2(F.relu(self.conv2(x)))
        x = self.bn3(F.relu(self.conv3(x)))
        x = self.pool(x)
        x = self.bn4(F.relu(self.conv4(x)))
        x = self.bn5(F.relu(self.conv5(x)))
```

```
# print(x.shape)
x = x.view(-1, 4 * 4 * 32)
x = F.relu(self.fc1(x))
x = self.dropout(x)
x = self.fc2(x)
return x
```

Обучим:

In [16]:

```
net = StrongConvNet()
train(net, learning_rate=0.001, epoch_num=10)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:12: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
if sys.path[0] == '':
```

```
/usr/local/lib/python3.7/dist-packages/torch/optim/lr_scheduler.py:134: UserWarning: Dete
cted call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later,
you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step(
)`. Failure to do this will result in PyTorch skipping the first value of the learning r
ate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate
"https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:17: TqdmDeprecationWarning:
This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`
```

```
[1, 2000] loss: 1.858
[1, 4000] loss: 1.665
[1, 6000] loss: 1.570
[1, 8000] loss: 1.522
[1, 10000] loss: 1.473
[1, 12000] loss: 1.398
```

```
[2, 2000] loss: 1.300
[2, 4000] loss: 1.293
[2, 6000] loss: 1.291
[2, 8000] loss: 1.292
[2, 10000] loss: 1.263
[2, 12000] loss: 1.233
```

```
[3, 2000] loss: 1.156
[3, 4000] loss: 1.149
[3, 6000] loss: 1.156
[3, 8000] loss: 1.148
[3, 10000] loss: 1.144
[3, 12000] loss: 1.148
```

```
[4, 2000] loss: 1.049
[4, 4000] loss: 1.031
[4, 6000] loss: 1.073
[4, 8000] loss: 1.057
[4, 10000] loss: 1.026
[4, 12000] loss: 1.039
```

```
[5, 2000] loss: 0.974
[5, 4000] loss: 0.984
[5, 6000] loss: 0.998
[5, 8000] loss: 0.966
[5, 10000] loss: 0.951
[5, 12000] loss: 0.955
```

```
[6, 2000] loss: 0.905
[6, 4000] loss: 0.910
[6, 6000] loss: 0.907
```

```
[6, 8000] loss: 0.909
[6, 10000] loss: 0.898
[6, 12000] loss: 0.913
```

```
[7, 2000] loss: 0.852
[7, 4000] loss: 0.838
[7, 6000] loss: 0.847
[7, 8000] loss: 0.861
[7, 10000] loss: 0.830
[7, 12000] loss: 0.860
```

```
[8, 2000] loss: 0.807
[8, 4000] loss: 0.825
[8, 6000] loss: 0.827
[8, 8000] loss: 0.797
[8, 10000] loss: 0.819
[8, 12000] loss: 0.827
```

```
[9, 2000] loss: 0.775
[9, 4000] loss: 0.797
[9, 6000] loss: 0.779
[9, 8000] loss: 0.781
[9, 10000] loss: 0.802
[9, 12000] loss: 0.791
```

```
[10, 2000] loss: 0.782
[10, 4000] loss: 0.793
[10, 6000] loss: 0.773
[10, 8000] loss: 0.795
[10, 10000] loss: 0.765
[10, 12000] loss: 0.764
```

Обучение закончено

In [17]:

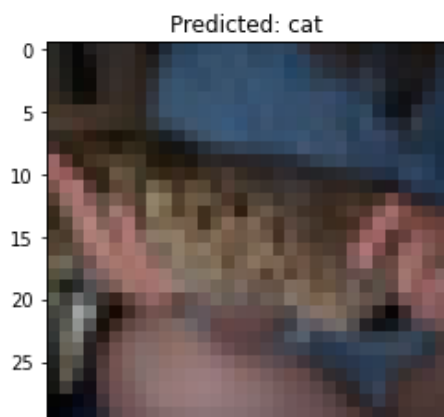
```
check_accuracy(net)
```


```
Accuracy of plane : 67 %
Accuracy of car : 77 %
Accuracy of bird : 50 %
Accuracy of cat : 48 %
Accuracy of deer : 60 %
Accuracy of dog : 51 %
Accuracy of frog : 72 %
Accuracy of horse : 69 %
Accuracy of ship : 75 %
Accuracy of truck : 73 %
Avg accuracy 64 %
```

Посмотрим визуально на работу нейросети:

In []:

```
i = np.random.randint(low=0, high=10000)
visualize_result(i)
```





Лучшая архитектура

Итого: наилучший результат **75%** точности в среднем, максимальной точности в **85%** для отдельного класса и минимальной точности в **57%** для отдельного класса. Результат был достигнут при **3x** сверточных слоях с увеличенным числом каналов и меньшим ядром свертки для получения большего пространства признаков (что приводит также к сильному увеличению времени обучения), с **avg** пулингом, **ReLU** в качестве функции активации и **3**мя линейными полносвязными слоями в качестве классификатора.

Даже обучив более глубокую и прокаченную (**BatchNorm**, **Dropout**) нейросеть на этих данных мы видим, что качество нас всё ещё не устраивает, в реальной жизни необходимо ошибаться не больше, чем на **5%**, а часто и это уже много. Как же быть, ведь свёрточные нейросети должны хорошо классифицировать изображения?

К сожалению, обучение нейросети с нуля на не очень большой выборке (а здесь она именно такая) часто приводит к переобучению, что плохо сказывается на тестовом качестве.

Для того, чтобы получить более качественную модель, часто дообучают сильную нейросеть, обученную на **ImageNet**, то есть используют технику **Transfer Learning**. О ней речь пойдёт далее в нашем курсе.

Полезные ссылки

1). Примеры написания нейросетей на **PyTorch** (официальные tutorиалы) (на английском):

https://pytorch.org/tutorials/beginner/pytorch_with_examples.html#examples

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

2). Курс Стэнфорда: <http://cs231n.github.io/>

3). Практически исчерпывающая информация по основам свёрточных нейросетей (из **cs231n**) (на английском):

<http://cs231n.github.io/convolutional-networks/>

<http://cs231n.github.io/understanding-cnn/>

<http://cs231n.github.io/transfer-learning/>

4). Видео о **Computer Vision** от Andrej Karpathy: <https://www.youtube.com/watch?v=u6aEYuemt0M>